

# A Note on Run Times and Scaling

David Dobor

January 29, 2016

We argue why a quadratic run-time of an algorithm is much too slow. The reason for this is that  $O(n^2)$  algorithms do not scale: as computers get faster and bigger, quadratic time algorithms actually get slower. We explain here what that means.

## What's wrong with $O(n^2)$ ?

Many of the algorithms we considered so far, like insertion sort or selection sort, had  $O(n^2)$  worst case running time. Large problems simply can not be solved by procedures that are this slow.

A rough standard for now is that people have computers that run billions of operations per second and have billions of entries in main memory. For example, I'm typing this on a 2.2 GHz device with 8 GBytes of main memory.

This means that you can touch everything in main memory on the order of about a second.

It is kind of an amazing fact that this rough standard has held for longer than over past half-a-century. The computers get bigger but they also get faster. To touch everything in memory is going to take a few seconds. This was true when computers had a few thousand words of memory, and it is true now when they have billions or more. So let's accept that for a moment as what computers are like.

What that means is that with that huge memory we can address huge problems. So we can have billions of objects and hope to sort them using bubble-sort, for example.

But there is a problem here. A  $O(n^2)$  algorithm taking as input  $n = 10^9$  entries will perform about  $(10^9)^2 = 10^{18}$  operations. This, on a computer that processes  $10^9$  entries per second, will take  $10^{18}/10^9 = 10^9$  seconds. Well, this works out to be over *30 years of computing time*. Obviously, it is not practical to address such a problem on today's computer.

NOW CONSIDER a computer that is 10 times as fast on which you can address a problem that is 10 times as big and which you will probably get to play with in not so distant future.




Figure 1: A modern processor.




Figure 2: A  $O(n^2)$  algorithm scales horribly.

The same algorithm will take  $(10 \text{ billion})^2 = (10^{10})^2 = 10^{20}$  operations on a device with  $10^{10}$  memory entries, all of which can be accessed in about a second. How long would that take?  $10^{20}/10^{10} = 10^{10}$  seconds. Well, this works out to be over *300 years of computing time*.

This is exactly the kind of situation you would want to try to avoid by developing algorithms that perform better.